

Apache Lucene – Block-Max Wand Algorithm

Apache Lucene performance improvements using block-max indexes and block-max WAND. As we all know Apache Lucene is open source high-performance, full-featured text search engine library.

In this document, I'm going to cover **BLOCK-MAX WAND ALGORITHM** that is used as a part of Boolean search queries in Apache Lucene.

Before we jump on to the basics of the algorithm, a small brief about the query processing techniques that are in use to tackle the longer processing times of queries.

Early termination is one important technique that addresses the above problem. We say that a query processing algorithm is exhaustive if it fully evaluates all documents that satisfy the Boolean filter condition. Any non-exhaustive algorithm is considered to use early termination (ET). There are four ways in which early termination often happens:

- **Stop early:** In this case, the postings are usually arranged such that the most promising documents appear early. Then we stop the traversal of the index as soon as we (may) have the top-k results. Well-known examples are the TA, FA, and NRA algorithms of Fagin.
- **Skip within lists:** When the postings in each list are sorted by docIDs, the promising documents are spread out throughout the inverted lists, and thus the standard intuition for "stop early" does not apply. There are few published works on early termination techniques under this scenario. **An exception is the WAND algorithm**, which uses a smart pointer movement technique to skip many documents that would be evaluated by an exhaustive algorithm.
- **Omit lists:** One or more lists for the query terms are completely ignored, if they do not affect the final results by much.
- **Score only partially:** We partially evaluate a document by computing only some term scores, or by computing approximate scores. When we find that the document cannot be in the top results, we stop evaluation.

The two index traversal techniques that we use in IR realm are:

Document-At-A-Time (DAAT): In DAAT query processing, each list has a pointer that points to a "current" posting in the list. All the pointers move forward in parallel as the query is being processed.

Term-At-A-Time (TAAT): In TAAT query processing, we first access one term, or one layer from one term, and then move to the next term, or the next layer from the same term or a different term. We use a temporary data structure to keep track of currently active top-k candidates.

WAND algorithm mainly uses DAAT as it consumes less space and uses less complex data structures. The algorithm primarily makes use of the global score defined on inverted lists and is less efficient when it comes to skipping the documents.

Apache Lucene – Block-Max Wand Algorithm

Block Max Wand Algorithm is an extension to WAND algorithm, where we compute impact score on block postings. A posting can have information about a document such as docid, term frequency, page rank score, similarity function used etc. We are splitting the posting into blocks of size 64 or 128 postings and calculate its impact score. The reason we are splitting the postings into multiple small blocks, is the average impact score of a inverted list may not accurately represent the documents present in it. Hence, by splitting into smaller chunks we can estimate the impact score more accurately.

Below is the algorithm for the same:

```
Initialize();
repeat
    /* sort the lists by current docIDs */
    Sort(lists);
    /* same "pivoting" as in WAND using the max
    impact for the whole lists, use p to denote
    the pivot */
    p = Pivoting(lists,  $\theta$ );
    d = lists[p]  $\rightarrow$  curDoc;
    if (d == MAXDOC) then
        | break;
    end
    for i = 0 ... p + 1 do
        | NextShallow(d, list(i));
    end
    flag = CheckBlockMax( $\theta$ , p);
    if (flag == true) then
        if ( lists[0]  $\rightarrow$  curDoc == d ) then
            | EvaluatePartial(d , p);
            | Move all pointers from lists[0] to lists[p] by calling
            | Next(list, d + 1)
        end
        else
            | Choose one list from the lists before lists[p] with the
            | largest IDF, move it by calling Next(list, d + 1)
        end
    end
    end
    else
        | d' = GetNewCandidate();
        | Choose one list from the lists before and including lists[p]
        | with the largest IDF, move it by calling Next(list, d')
    end
until Stop;
```

Apache Lucene – Block-Max Wand Algorithm

While we split the postings into blocks, we store the block boundaries in the inverted index table. A global score is an additional information about a document such as its length, pagerank etc. We use the global score information. We use this global score and the doc id to calculate the candidate pivot. Once we have the pivot, we use the `NextShallow` method to identify whether the pivot is indeed a correct one. This method reduces the number of blocks that we need to decompress there by improving the performance of the search query.

In apache Lucene, a **scorer** is an abstract class defined to expose the below functionalities:

- o Expose an `iterator()` over documents matching a query in increasing order of doc Id.
- o Compute Document scores using a given Similarity implementation.

DocIdSetIterator class defines methods to iterate over a set of non-decreasing doc ids. The max doc ids value is set to 2147483647 in order to be used as a sentinel object.

Classes `BlockMaxConjunctionScorer`, `WANDScorer` and `DisjunctionSumScorer` implements the **scorer** interface and make use of the **DocIdSetIterator** in order to access the documents.

```
// Technically speaking, WANDScorer should be able to handle the following 3 conditions now
// 1. Any ScoreMode (with scoring or not)
// 2. Any minCompetitiveScore ( >= 0 )
// 3. Any minShouldMatch ( >= 0 )
//
// However, as WANDScorer uses more complex algorithm and data structure, we would like to
// still use DisjunctionSumScorer to handle exhaustive pure disjunctions, which may be faster
if (scoreMode == ScoreMode.TOP_SCORES || minShouldMatch > 1) {
    return new WANDScorer(weight, optionalScorers, minShouldMatch, scoreMode);
} else {
    return new DisjunctionSumScorer(weight, optionalScorers, scoreMode);
}
```

Disjunctive queries(queries with OR operator) make use of **WANDScorer** or **DisjunctionSumScorer** whereas Conjunction queries(queries with AND operator) make use of **BlockMaxConjunctionScorer**.

Apache Lucene – Block-Max Wand Algorithm

Advance to the block of documents that contains `target` in order to get scoring information about this block. This method is implicitly called by `DocIdSetIterator.advance(int)` and `DocIdSetIterator.nextDoc()` on the returned doc ID. Calling this method doesn't modify the current `DocIdSetIterator.docID()`. It returns a number that is greater than or equal to all documents contained in the current block, but less than any doc IDS of the next block. `target` must be $\geq \text{docID}()$ as well as all targets that have been passed to `advanceShallow(int)` so far.

```
public int advanceShallow(int target) throws IOException {  
    return DocIdSetIterator.NO_MORE_DOCS;  
}
```

As described in the above algorithm, the shallow, method has been implemented in Lucene but with a small change, instead of storing the block score value in the index table, Lucene is storing the term frequency and document id, using the two we can calculate the block index value.

```
public interface ImpactsSource {
```

Shallow-advance to `target`. This is cheaper than calling `DocIdSetIterator.advance(int)` and allows further calls to `getImpacts()` to ignore doc IDs that are less than `target` in order to get more precise information about impacts. This method may not be called on targets that are less than the current `DocIdSetIterator.docID()`. After this method has been called, `DocIdSetIterator.nextDoc()` may not be called if the current doc ID is less than `target - 1` and `DocIdSetIterator.advance(int)` may not be called on targets that are less than `target`.

```
void advanceShallow(int target) throws IOException;
```

Get information about upcoming impacts for doc ids that are greater than or equal to the maximum of `DocIdSetIterator.docID()` and the last target that was passed to `advanceShallow(int)`. This method may not be called on an unpositioned iterator on which `advanceShallow(int)` has never been called. NOTE: advancing this iterator may invalidate the returned impacts, so they should not be used after the iterator has been advanced.

```
Impacts getImpacts() throws IOException;
```

```
}
```

The score functions are making use of the `Lucene84PostingsFormat` class that is responsible for compressing and decompressing the index data. It has also implemented functions such as `advanceShallow(int target)` and `getImpacts()` to further reduce the burden on the `DocIdSetIterator`.

With the performance and optimization benefits, we also must compromise on some of the other stuff, such as accurate hit counts. Using this we cannot get the accurate hit counts and in order to get that we will have hit on performance.

The implementation above made term queries run between 3x and 7x faster, conjunctions between 3% and 7x faster and disjunctions between -8% (slightly slower) and 15x faster when running Lucene's benchmark suite.

Apache Lucene – Block-Max Wand Algorithm

References:

1. <http://engineering.nyu.edu/~suel/papers/bmw.pdf>
2. <https://lucene.apache.org/core/>
3. <https://www.elastic.co/blog/faster-retrieval-of-top-hits-in-elasticsearch-with-block-max-wand>